



University of Minho  
School of Engineering



# Aprendizagem Profunda

## CVAE, GANs, FCN

APP @ MEI/1º ano – 2º Semestre

Victor Alves

Part VIII

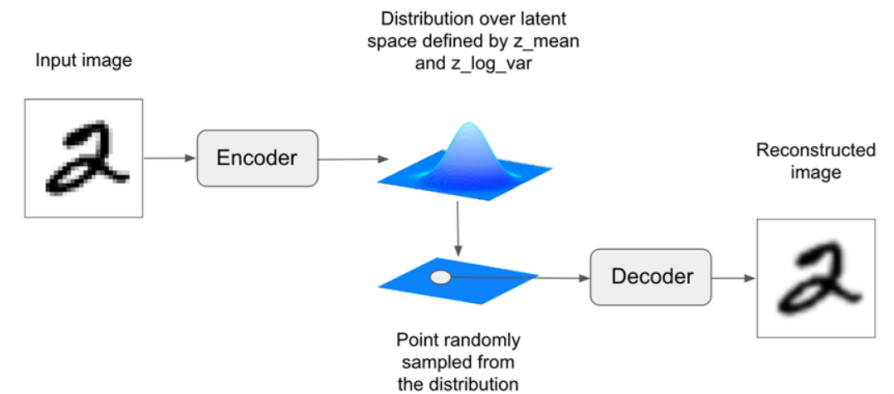

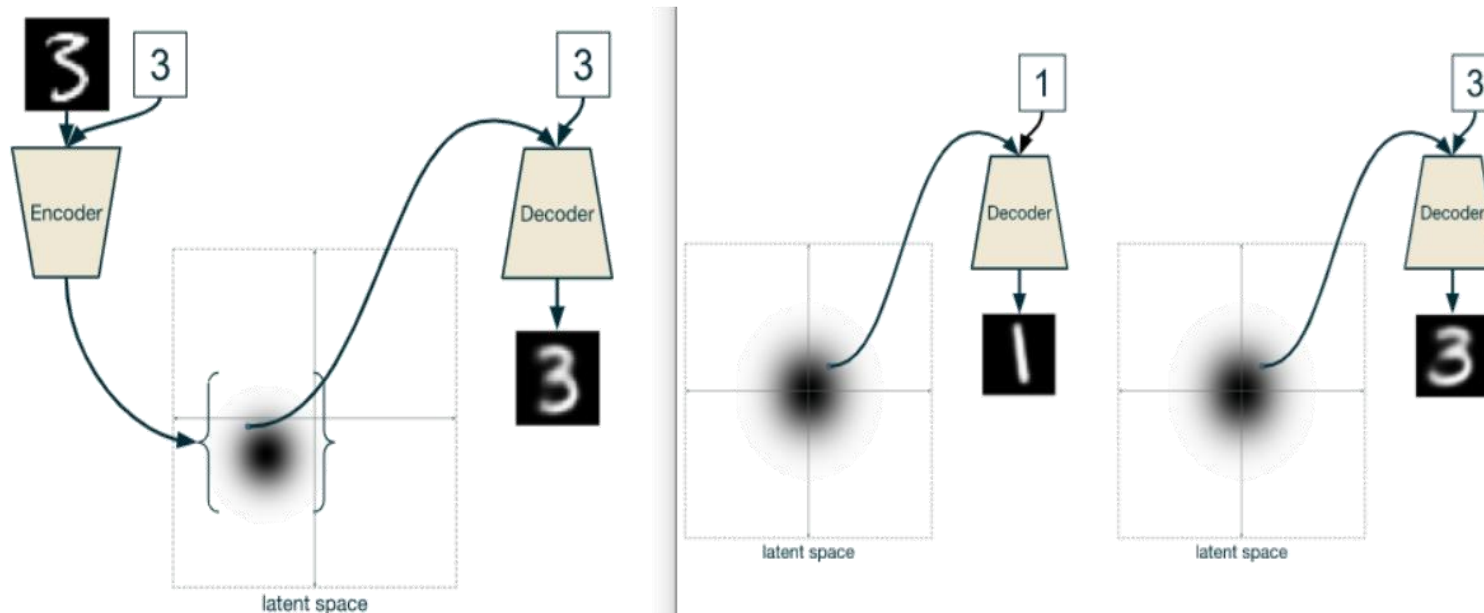
# Contents

---

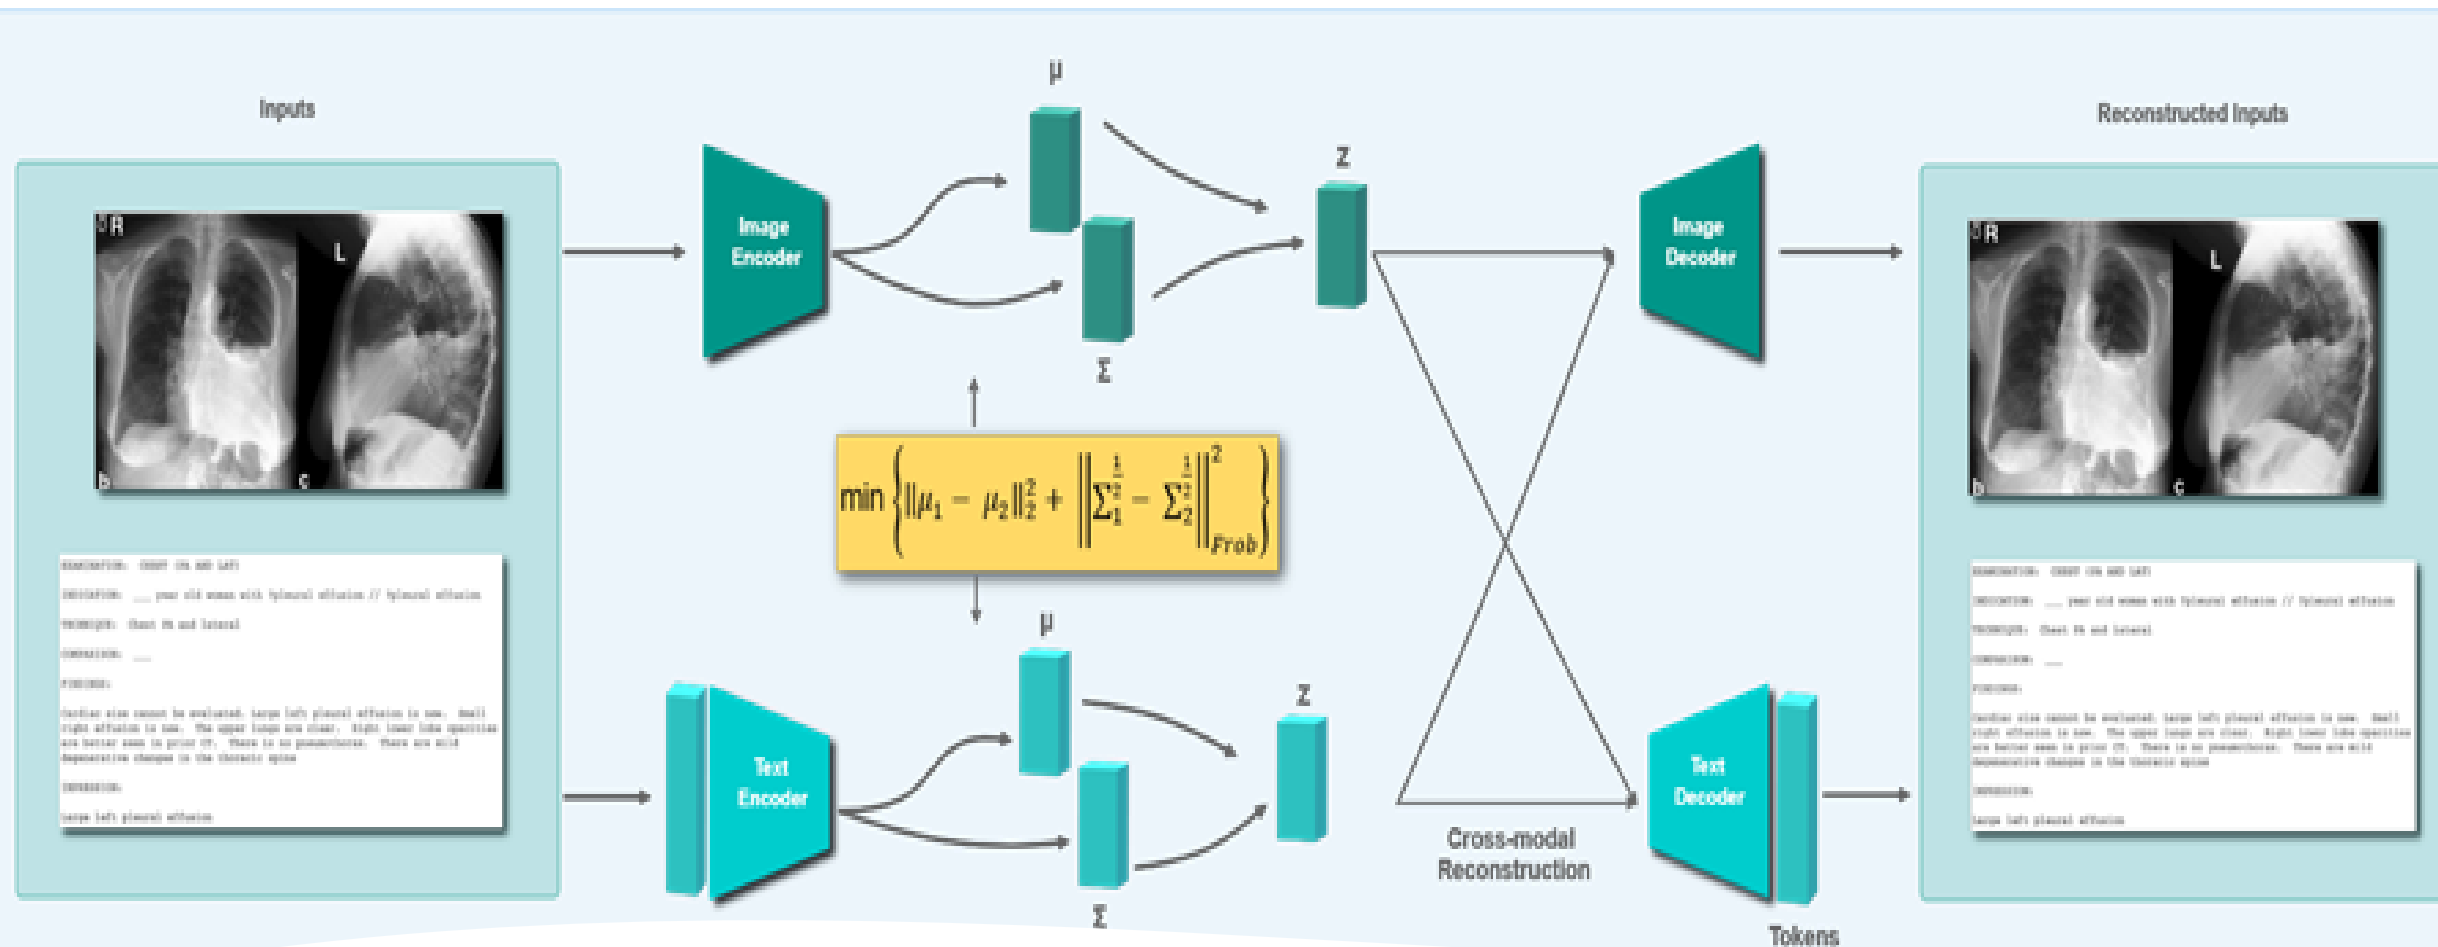
- CVAE
- GANs
- FCN

# Laboratory Classes

- Conditional Variational Autoencoder with MNIST dataset
  - 11\_pytorch\_CVAE\_MLP\_training\_MNIST.ipynb
  - 12\_pytorch\_CVAE\_MLP\_generate\_MNIST.ipynb



# Zero Shot Learning

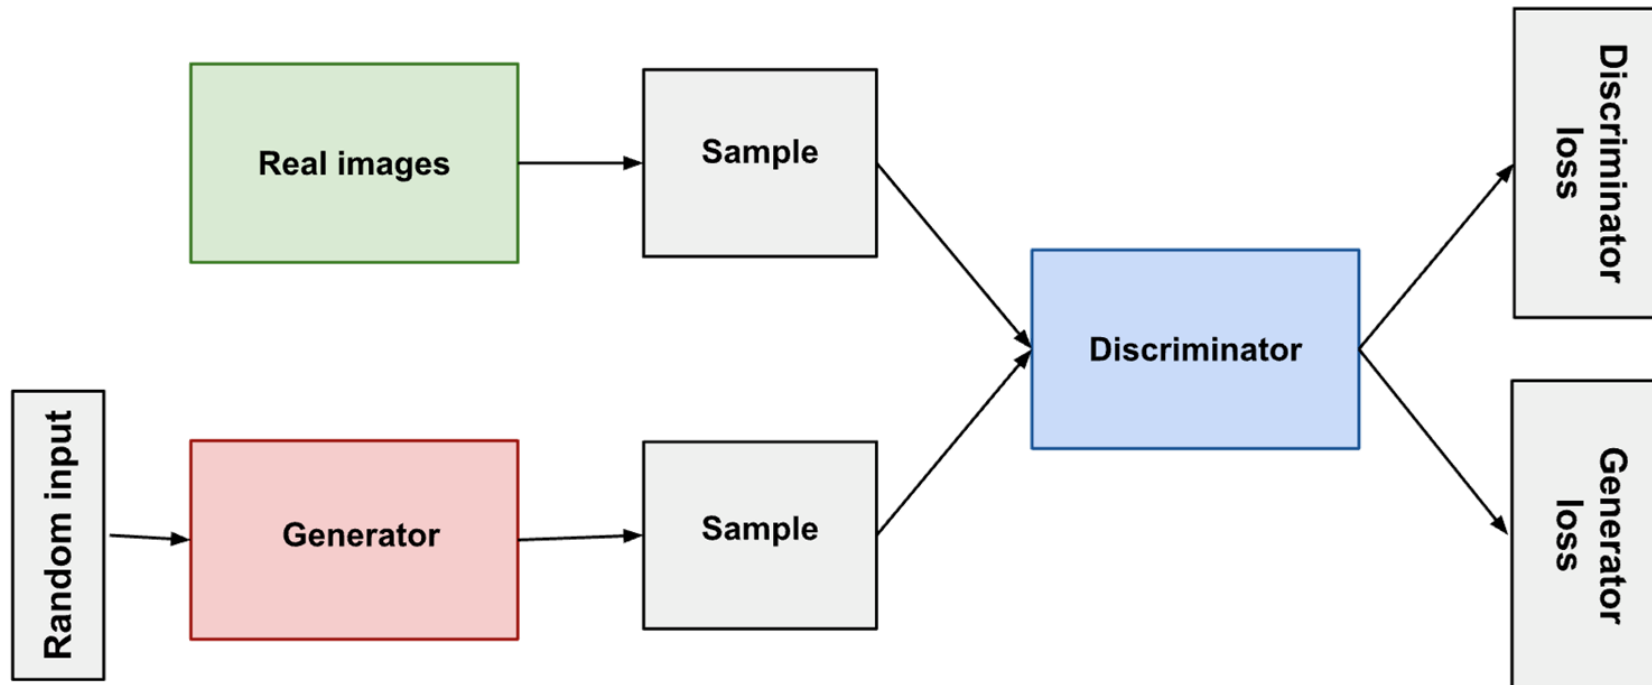


# GAN - Generative Adversarial Networks

- Introduced in 2014 by Ian Goodfellow et al.
- They have been used essentially to **generate new images that are statistically similar to the original ones** (among other types of data)
- GAN includes **two models**: one that generates new examples and another that tries to discriminate whether examples are new or original (used in training), trained to optimise opposite (adversarial) functions (loss)
- An intuitive example would be a forger of paintings (or banknotes) showing his work to an art expert; as the forger gets better at creating better forgeries, the expert gets better at recognizing them

# GAN - Generative Adversarial Networks

- **Generator** - takes as input a random vector (point generated in latent space), decoding it into a new image
- **Discriminator** (adversary) - receives as input an image (real or generated) and classifies it as original (from the training data set) or one generated by the Generator



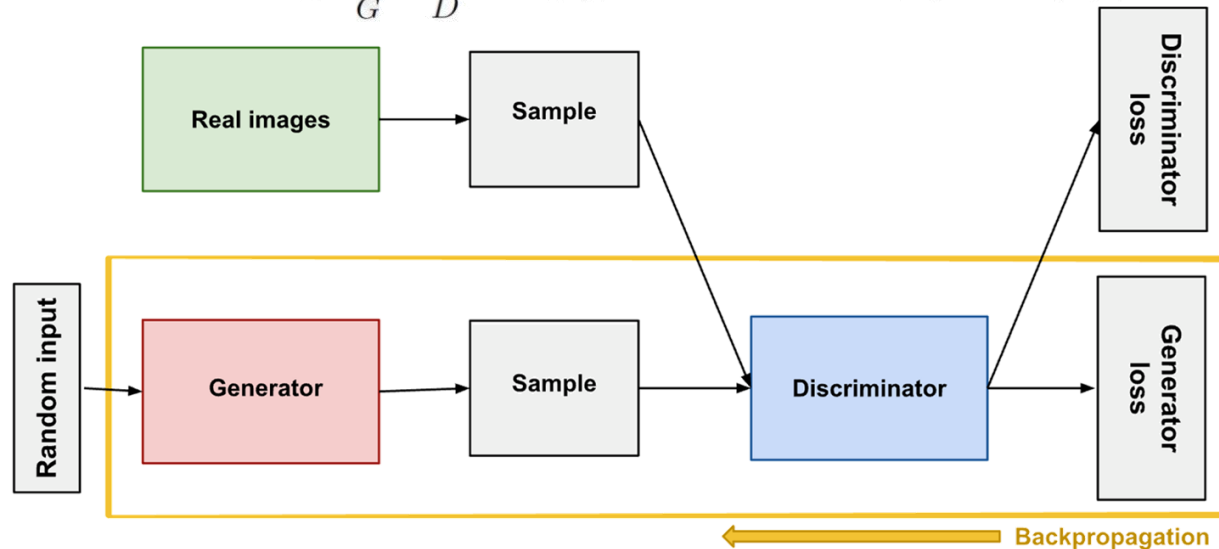
# GAN - Generative Adversarial Networks

- **Generator** : receives a random value (noise) which it transforms into a new example (e.g. image); seeks to **minimise** the probability of the discriminator distinguishing the data it generates as being false.

$$\arg \min_G \mathbb{E}_{\mathbf{z}, \mathbf{x}} [ \log D(G(\mathbf{z})) + \log (1 - D(\mathbf{x})) ]$$

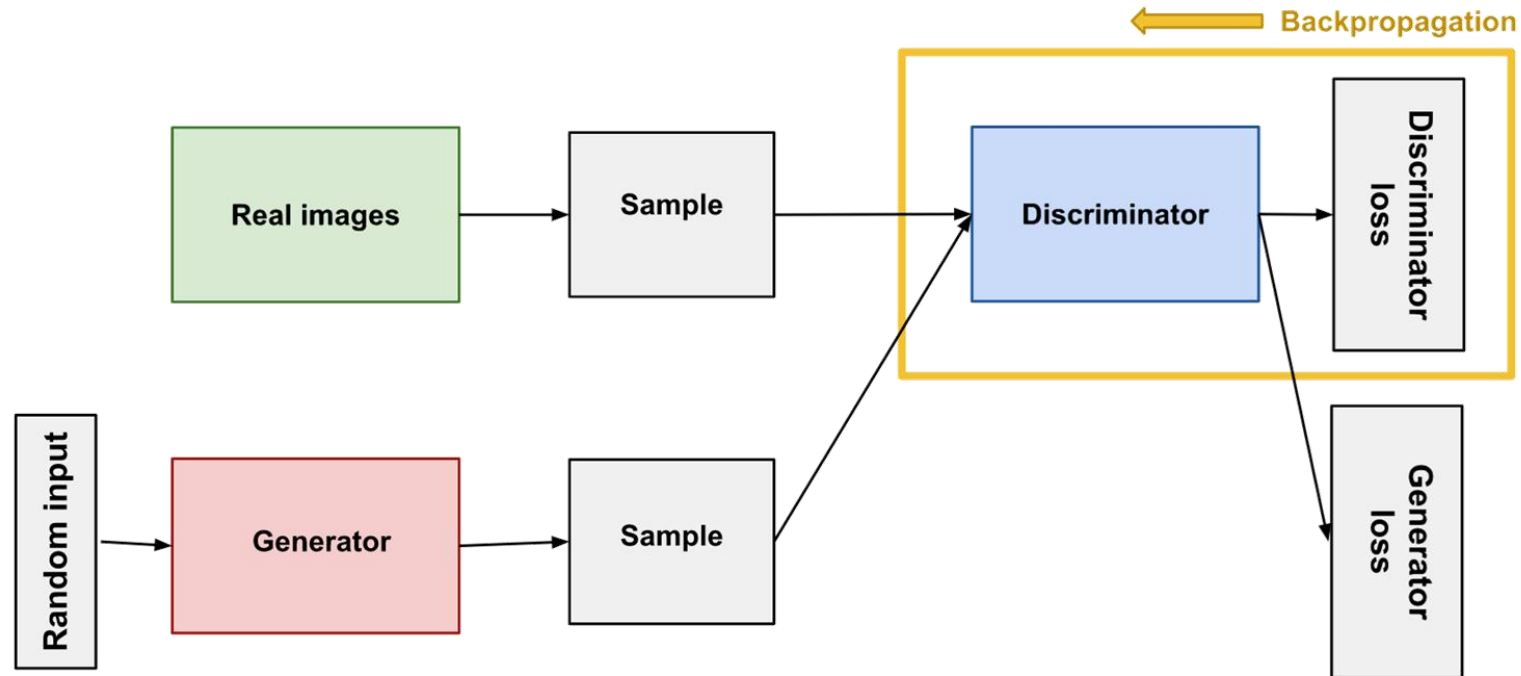
- The ultimate end of the generator is to generate false data that deceives the discriminator:

$$\arg \min_G \max_D \mathbb{E}_{\mathbf{z}, \mathbf{x}} [ \log D(G(\mathbf{z})) + \log (1 - D(\mathbf{x})) ]$$



# GAN - Generative Adversarial Networks

- **Discriminator:** maximizes the probability of identifying false data as false.



$$\arg \max_{\underline{D}} \mathbb{E}_{\mathbf{z}, \mathbf{x}} \left[ \underbrace{\log D(G(\mathbf{z}))}_{\text{Fake}} + \underbrace{\log (1 - D(\mathbf{x}))}_{\text{Real}} \right]$$



# GAN - Generative Adversarial Networks

**Training** - must alternate between Generator and Discriminator training, one or more epochs each (while training one model, the other remains unchanged)

## **Generator training:**

- Generate fake images
- Compute the discriminator loss on fake images
- Perform backpropagation + an optimization step to update the generator's weights

## **Discriminator training:**

- Compute the discriminator loss on real, training images
- Generate fake images
- Compute the discriminator loss on fake, generated images
- Add up the real and fake loss
- Perform backpropagation + an optimization step to update the discriminator's weights

# GAN - Generative Adversarial Networks

- **Convergence** is difficult to identify because the loss functions are opposite (adversarial) and therefore there is no absolute improvement
- Training of GANs is quite difficult in practice - several "**tricks**" needed, e.g:
  - Use tanh as the last activation in the generator, instead of sigmoid;
  - Sample latent space points using a normal distribution (Gaussian distribution) instead of a uniform distribution.
  - Introduce randomness by using dropout in the discriminator or by adding random noise to the discriminator labels.

# Other GANs / applications

- <https://developers.google.com/machine-learning/gan/applications>
- **Progressive GANs** – higher resolution images
- **Conditional GANs** – create constraints for the images to be generated
- Image to image translation
- Text to image synthesis
- Text to speech
- **Cycle GANs** – image processing



# Laboratory Classes


---

- Generative Adversarial Networks with medMNIST dataset
  - 13\_pytorch\_GAN\_medNIST.ipynb


# FCN - Fully Convolutional Network




Object Detection




Semantic Segmentation



Instance Segmentation




Input Image



Segmented Image

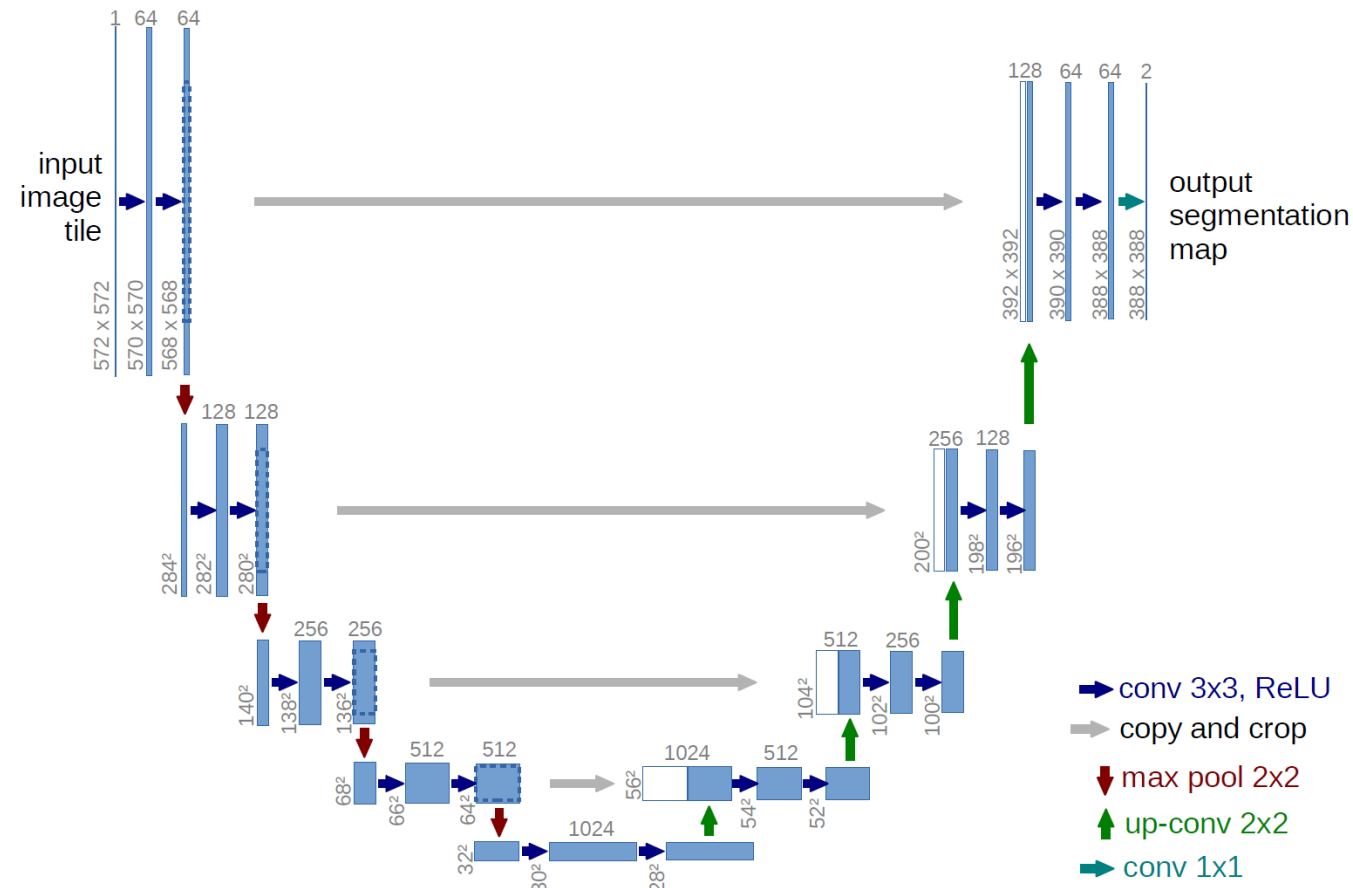
# FCN - Fully Convolutional Network

## SegNet



# FCN - Fully Convolutional Network

## • U-net



# Laboratory Classes

- Segmentation Exercise
  - 16.1\_pytorch\_exercicio\_segmentacao.ipynb

In this exercise we will segment the left ventricle of the heart in relatively small images using neural nets.

The code for a segmentation net and its training is presented. The network is not very good, so **the exercise is to improve the quality of segmentation by improving the network and/or the training scheme, including data loading efficiency and data augmentation.**

The data used here are derived from [Sunnybrook Cardiac Dataset cardiac](#) MR images, filtered to contain only segmentations of the left ventricular myocardium and reduced in XY dimensions.

Data extracted from:

[https://github.com/ericspod/VPHSummerSchool2019/raw/master/scd\\_lvsegs.npz](https://github.com/ericspod/VPHSummerSchool2019/raw/master/scd_lvsegs.npz)



# Occlusion

## Sensitivity to Occlusion

One method for trying to visualise why the net made a particular prediction is occlusion sensitivity. We occlude part of the image and see how the probability of a particular prediction changes. We then iterate over the image, moving the occluded part as we go, and in doing so build a sensitivity map detailing which areas were the most important in the decision making.

