



University of Minho
School of Engineering



Aprendizagem Profunda

Deep Neural Network – CNN Multiclass

APP @ MEI/4º ano – 2º Semestre

Victor Alves, Miguel Rocha



Hands On

CNN for multiclass image classification

MNIST (Modified National Institute of Standards and Technology) Dataset

It is considered the "hello world" dataset of *computer vision*.

- Dataset of manually written digit images
- Inputs: 28x28 pixels images
- Output: class representing the digit (10 classes, digits 0-9)
- 70k images of which 60k are for training and 10k for testing
- 2 attributes: the image id and its label

We will use a convolutional neural network to classify the digit in each 28x28 image.

0. Prepare the setup

Install pytorch (if needed)

Imports

Constants

Device management (optional)

0. Prepare the setup

```
device = torch.device("cuda" if
torch.cuda.is_available() else "cpu")

def get_default_device():
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)
```

```
class DeviceDataLoader():
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self.dl)
```

```
device = get_default_device()
print(device)
```

1. Prepare the data

```
train_transform = Compose(  
    [ToTensor(),  
     Normalize(mean=(0.1307,), std=(0.3081,))  
    ]  
test_transform = Compose(  
    [ToTensor(),  
     Normalize(mean=(0.1307,), std=(0.3081,))  
    ])
```

1. Prepare the data

```
class CSVDataset(Dataset):

    def __init__(self, path, transform=None):
        self.transform = transform
        df_set = pd.read_csv(path, header=0)
        self.x = df_set.values[:, 1:]
        self.y = df_set.values[:, 0]
        self.x = self.x.astype('float32')
        self.y = self.y.astype('long')

    def __len__(self):
        return len(self.x)
```

```
    def __getitem__(self, idx):
        label = self.y[idx]
        image = self.x[idx]
        if self.transform is not None:
            image = self.transform(image)
        return image, label

    def get_TensorDataset(self):
        x = self.x.reshape(len(self.x), 1, 28, 28)
        xmax, xmin = x.max(), x.min()
        x = (x - xmin)/(xmax - xmin)
        x = torch.from_numpy(np.array(x)).float()
        y = torch.from_numpy(np.array(self.y)).type(
            torch.LongTensor)
        cases = torch.utils.data.TensorDataset(x,y)
        return cases
```

1. Prepare the data

```
def prepare_data_loaders(path_train, path_test):
    dataset_train = CSVDataset(path_train, transform=train_transform)
    dataset_test = CSVDataset(path_test, transform=test_transform)
    train = dataset_train.get_TensorDataset()
    train_size = int(0.8 * len(train))
    val_size = len(train) - train_size
    train, validation = random_split(train, [train_size, val_size], generator=torch.Generator().manual_seed(42))
    test = dataset_test.get_TensorDataset()
    train_dl = DataLoader(train, batch_size=BATCH_SIZE, shuffle=True)
    val_dl = DataLoader(validation, batch_size=BATCH_SIZE, shuffle=True)
    test_dl = DataLoader(test, batch_size=BATCH_SIZE, shuffle=True)
    train_dl_all = DataLoader(train, batch_size=len(train), shuffle=False)
    val_dl_all = DataLoader(validation, batch_size=len(validation), shuffle=True)
    test_dl_all = DataLoader(test, batch_size=len(test), shuffle=False)
    return train_dl, val_dl, test_dl, train_dl_all, val_dl_all, test_dl_all
```

```
train_dl, val_dl, test_dl, train_dl_all, val_dl_all, test_dl_all = prepare_data_loaders(PATH_TRAIN, PATH_TEST)
```


1.1 Visualize the data

```
def output_label(label,mapping='label'):
    if mapping == 'ext':
        output_mapping = { 0:"zero", 1:"um", 2:"dois", 3:"tres", 4:"quatro", 5:"cinco", 6:"seis",
7:"sete",
                        8:"oito", 9:"nove" }
    elif mapping == 'ext2':
        output_mapping = { "0":"zero", "1":"um", "2":"dois", "3":"tres", "4":"quatro", "5":"cinco",
                        "6":"seis", "7":"sete", "8":"oito", "9":"nove" }
    else:
        output_mapping = { 0: "0", 1: "1", 2: "2", 3: "3", 4: "4", 5: "5", 6: "6", 7: "7", 8: "8", 9: "9"}
    input = (label.item() if type(label) == torch.Tensor else label)
    return output_mapping[input]
```

1.1 Visualize the data

```
from IPython.display import display
def visualize_data(path):
    df = pd.read_csv(path, header=0)
    display(df)
```

```
def visualize_dataset(train_dl, test_dl):
    print(f"Quantidade de casos de Treino:{len(train_dl.dataset)}")
    print(f"Quantidade de casos de Validação:{len(val_dl.dataset)}")
    print(f"Quantidade de casos de Teste:{len(test_dl.dataset)}")
    x, y = next(iter(train_dl))
    print(f"Shape tensor batch casos treino, input: {x.shape}, output: {y.shape}")
    x, y = next(iter(val_dl))
    print(f"Shape tensor batch casos validação, input: {x.shape}, output: {y.shape}")
    x, y = next(iter(test_dl))
    print(f"Shape tensor batch casos test, input: {x.shape}, output: {y.shape}")
    print(f'Valor máximo:{torch.max(x)} Valor mínimo:{torch.min(x)}')
    x=x.detach().numpy()
    print(f'Valor máximo:{np.max(x)} Valor mínimo:{np.min(x)}')
    print(y)
```

```
visualize_data(PATH_TRAIN)
```

1.1 Visualize the data

	label	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	...	28x19	28x20	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
0	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
...
59995	8	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59996	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59997	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59998	6	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
59999	8	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

60000 rows × 785 columns

Quantidade de casos de Treino:48000

Quantidade de casos de Validação:12000

Quantidade de casos de Teste:10000

Shape tensor batch casos treino, input: torch.Size([32, 1, 28, 28]), output: torch.Size([32])

Shape tensor batch casos validação, input: torch.Size([32, 1, 28, 28]), output: torch.Size([32])

Shape tensor batch casos test, input: torch.Size([32, 1, 28, 28]), output: torch.Size([32])

Valor máximo:1.0 Valor mínimo:0.0

Valor máximo:1.0 Valor mínimo:0.0

tensor([7, 1, 7, 4, 3, 4, 2, 8, 7, 4, 2, 5, 8, 3, 2, 5, 1, 1, 5, 4, 3, 2, 1, 2,
3, 9, 5, 6, 7, 7, 3, 9])

1.1 Visualize the data

```
def visualize_mnist_images(dl):
    i, (inputs, targets) = next(enumerate(dl))
    print(inputs.shape)
    plt.figure(figsize=(8,8))
    for i in range(25):
        plt.subplot(5, 5, i+1)
        plt.axis('off')
        plt.grid(b=None)
        plt.imshow(inputs[i][0], cmap='gray')
    plt.show()


visualize_mnist_images(train_dl)
```

1.1 Visualize the data

```
def visualize_mnist_images(dl):
    i, (inputs, targets) = next(enumerate(dl))
    print(inputs.shape)
    plt.figure(figsize=(8,8))
    for i in range(25):
        plt.subplot(5, 5, i+1)
        plt.axis('off')
        plt.grid(b=None)
        plt.imshow(inputs[i][0], cmap='gray')
    plt.show()

visualize_mnist_images(train_dl)
```

torch.Size([32, 1, 28, 28])



1.2 Verify dataset balancing

Sequential Layer 1(Conv, ReLU, MaxPool)
Sequential Layer 2(Conv, ReLU, MaxPool)
Linear
ReLU
Linear
Softmax

2. Define the model (1)

```
class CNNModel_1(Module):
    def __init__(self):
        super(CNNModel_1, self).__init__()
        self.layer1 = Sequential(Conv2d(in_channels=1, out_channels=32, kernel_size=(3,3)),
                                   ReLU(),
                                   MaxPool2d(kernel_size=(2,2), stride=(2,2))
                                   )
        self.layer2 = Sequential(Conv2d(in_channels=32, out_channels=32, kernel_size=(3,3)),
                                   ReLU(),
                                   MaxPool2d(kernel_size=(2,2), stride=(2,2))
                                   )
        self.fc1 = Linear(in_features=5*5*32, out_features=100)
        kaiming_uniform_(self.fc1.weight, nonlinearity='relu')
        self.act1 = ReLU()
        self.fc2 = Linear(in_features=100, out_features=10)
        xavier_uniform_(self.fc2.weight)
        self.act2 = Softmax(dim=1)
```

2. Define the model (1)

Sequential Layer 1(Conv, ReLU, MaxPool)
Sequential Layer 2(Conv, ReLU, MaxPool)
Linear
ReLU
Linear
Softmax

```
def forward(self, x):  
    out = self.layer1(x)  
    out = self.layer2(out)  
    out = out.view(-1, 4*4*50)  
    out = self.fc1(out)  
    out = self.act1(out)  
    out = self.fc2(out)  
    out = self.act2(out)  
    return out
```

```
model = CNNModel_1()  
print(summary(model, input_size=(BATCH_SIZE, 1,28,28), verbose=0))  
model.to(device)
```

2. Define the model (1)

Sequential Layer 1(Conv, ReLU, MaxPool)
Sequential Layer 2(Conv, ReLU, MaxPool)
Linear
ReLU
Linear
Softmax

```
=====
Layer (type:depth-idx)                   Output Shape           Param #
=====
└─Sequential: 1-1                        [32, 32, 13, 13]      --
    └─Conv2d: 2-1                         [32, 32, 26, 26]      320
        └─ReLU: 2-2                       [32, 32, 26, 26]      --
            └─MaxPool2d: 2-3               [32, 32, 13, 13]      --
└─Sequential: 1-2                        [32, 32, 5, 5]        --
    └─Conv2d: 2-4                         [32, 32, 11, 11]      9,248
        └─ReLU: 2-5                       [32, 32, 11, 11]      --
            └─MaxPool2d: 2-6               [32, 32, 5, 5]        --
└─Linear: 1-3                            [32, 100]              80,100
└─ReLU: 1-4                             [32, 100]              --
└─Linear: 1-5                             [32, 10]               1,010
└─Softmax: 1-6                           [32, 10]               --
=====
Total params: 90,678
Trainable params: 90,678
Non-trainable params: 0
Total mult-adds (M): 54.24
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 6.56
Params size (MB): 0.36
Estimated Total Size (MB): 7.02
=====
```


3. Train the model

17

```
def train_model(h5_file, train_dl, val_dl, model, criterion, optimizer):
```

```
    liveloss = PlotLosses()
```

```
    for epoch in range(EPOCHS):
```

```
        logs = {}
```

```
        model.train()
```

```
        running_loss = 0.0
```

```
        running_corrects = 0.0
```

```
        for inputs, labels in train_dl:
```

```
            inputs = inputs.to(device)
```

```
            labels = labels.to(device)
```

```
            outputs = model(inputs)
```

```
            loss = criterion(outputs, labels)
```

```
            optimizer.zero_grad()
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
            running_loss += loss.detach() * inputs.size(0)
```

```
            _, preds = torch.max(outputs, 1)
```

```
            running_corrects += torch.sum(preds == labels.data)
```

```
    (...)
```

```
    (...)
```

```
    epoch_loss = running_loss / len(train_dl.dataset)
```

```
    epoch_acc = running_corrects.float()/len(train_dl.dataset)
```

```
    logs['loss'] = epoch_loss.item()
```

```
    logs['accuracy'] = epoch_acc.item()
```

```
    model.eval()
```

```
    running_loss = 0.0
```

```
    running_corrects = 0.0
```

```
    for inputs, labels in val_dl:
```

```
        inputs = inputs.to(device)
```

```
        labels = labels.to(device)
```

```
        outputs = model(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        running_loss += loss.detach() * inputs.size(0)
```

```
        _, preds = torch.max(outputs, 1)
```

```
        running_corrects += torch.sum(preds == labels.data)
```

```
    epoch_loss = running_loss / len(val_dl.dataset)
```

```
    epoch_acc = running_corrects.float() / len(val_dl.dataset)
```

```
    logs['val_loss'] = epoch_loss.item()
```

```
    logs['val_accuracy'] = epoch_acc.item()
```

```
    liveloss.update(logs)
```

```
    liveloss.send()
```

```
    torch.save(model, h5_file)
```

3. Train the model (1)

18

For model 1:

```
model = CNNModel_1()
print(summary(model, input_size=(BATCH_SIZE, 1,28,28), verbose=0))
model.to(device)

EPOCHS = 15
LEARNING_RATE = 0.001
criterion = CrossEntropyLoss()
optimizer = SGD(model.parameters(), lr=LEARNING_RATE)
starttime = time.perf_counter()
train_model('CNNModel_1.pth', train_dl, val_dl, model, criterion, optimizer)
endtime = time.perf_counter()
print(f"Tempo gasto: {endtime - starttime} segundos")
```

3. Train the model (1)


Sequential Layer 1(Conv, ReLU, MaxPool)
Sequential Layer 2(Conv, ReLU, MaxPool)
Linear
ReLU
Linear
Softmax

19

```
=====
Layer (type:depth-idx)      Output Shape      Param #
=====
|Sequential: 1-1            [32, 32, 13, 13]  --
|  └─Conv2d: 2-1            [32, 32, 26, 26]  320
|    └─ReLU: 2-2            [32, 32, 26, 26]  --
|      └─MaxPool2d: 2-3     [32, 32, 13, 13]  --
|Sequential: 1-2            [32, 32, 5, 5]    --
|  └─Conv2d: 2-4            [32, 32, 11, 11]  9,248
|    └─ReLU: 2-5            [32, 32, 11, 11]  --
|      └─MaxPool2d: 2-6     [32, 32, 5, 5]    --
|Linear: 1-3                [32, 100]         80,100
|ReLU: 1-4                  [32, 100]         --
|Linear: 1-5                [32, 10]          1,010
|Softmax: 1-6               [32, 10]          --
=====
Total params: 90,678
Trainable params: 90,678
Non-trainable params: 0
Total mult-adds (M): 54.24
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 6.56
Params size (MB): 0.36
Estimated Total Size (MB): 7.02
=====
```

3. Train the model (1)

20



```
accuracy
  training      (min: 0.147, max: 0.808, cur: 0.808)
  validation    (min: 0.197, max: 0.804, cur: 0.804)
Loss
  training      (min: 1.674, max: 2.299, cur: 1.674)
  validation    (min: 1.678, max: 2.296, cur: 1.678)
Tempo gasto: 2236.4488256331533 segundos
```

4. Evaluate the model

21

```
def evaluate_model(test_dl, model):
    predictions = list()
    actual_values = list()
    for inputs, labels in test_dl:
        inputs = inputs.to(device)
        labels = labels.to(device)
        yprev = model(inputs)
        yprev = yprev.detach().cpu().numpy()
        actual = labels.cpu().numpy()
        yprev = np.argmax(yprev, axis=1)
        actual = actual.reshape((len(actual), 1))
        yprev = yprev.reshape((len(yprev), 1))
        predictions.append(yprev)
        actual_values.append(actual)
    break
predictions, actual_values = np.vstack(predictions), np.vstack(actual_values)
return actual_values, predictions
```

4. Evaluate the model

22

```
def display_predictions(actual_values, predictions):
    acertou=0
    falhou = 0
    primeiros=0
    for r,p in zip(actual_values, predictions):
        if primeiros <20:
            print(f'real:{r} previsão:{p}')
            primeiros +=1
        if r==p: acertou+=1
        else: falhou+=1
    corrects = np.sum(predictions == actual_values)
    acc = corrects/len(test_dl.dataset)
    (...)
```

```
(...)
    acc = accuracy_score(actual_values,
predictions)
    print(f'Accuracy: {acc:0.3f}\n')
    print(f'acertou:{acertou} falhou:{falhou}')
    acc = accuracy_score(actual_values,
predictions)
    print(f'Accuracy: {acc:0.3f}\n')
    print(f'acertou:{acertou} falhou:{falhou}')
```

4. Evaluate the model

23

```
def display_confusion_matrix(cm,list_classes):  
    plt.figure(figsize = (16,8))  
    sns.heatmap(cm,annot=True,xticklabels=list_classes,yticklabels=list_classes, annot_kws={"size": 12},  
                fmt='g', linewidths=.5)  
  
    plt.ylabel('True label')  
    plt.xlabel('Predicted label')  
    plt.show()  
  
model= torch.load('CNNModel_1.pth')  
actual_values, predictions = evaluate_model(test_dl_all, model)  
display_predictions(actual_values, predictions )  
print(classification_report(actual_values, predictions))  
cr =classification_report(actual_values, predictions, output_dict=True)  
list_classes=[output_label(n,'ext2') for n in list(cr.keys())[0:10] ]  
cm = confusion_matrix(actual_values, predictions)  
print (cm)  
display_confusion_matrix(cm,list_classes)
```


4. Evaluate the model

24

real:[7] previsão:[7]		precision	recall	f1-score	support
real:[2] previsão:[2]					
real:[1] previsão:[1]	0	0.97	0.99	0.98	980
real:[0] previsão:[0]	1	0.98	0.99	0.99	1135
real:[4] previsão:[4]	2	0.97	0.96	0.97	1032
real:[1] previsão:[1]	3	0.98	0.98	0.98	1010
real:[4] previsão:[4]	4	0.98	0.98	0.98	982
real:[9] previsão:[9]	5	0.98	0.98	0.98	892
real:[5] previsão:[5]	6	0.99	0.97	0.98	958
real:[9] previsão:[9]	7	0.97	0.96	0.96	1028
real:[0] previsão:[0]	8	0.97	0.97	0.97	974
real:[6] previsão:[6]	9	0.97	0.96	0.97	1009

	accuracy				
real:[9] previsão:[9]				0.98	10000
real:[0] previsão:[0]	macro avg	0.98	0.98	0.98	10000
real:[1] previsão:[1]	weighted avg	0.98	0.98	0.98	10000

real:[9] previsão:[9]	[[969 0 2 0 0 0 3 2 4 0]
real:[7] previsão:[7]	[0 1128 1 2 1 1 1 0 1 0]
real:[3] previsão:[3]	[6 5 993 3 5 0 2 9 8 1]
real:[4] previsão:[4]	[0 0 2 993 0 5 0 6 2 2]
Accuracy: 0.975	[1 0 1 0 965 0 4 1 3 7]
	[3 1 0 4 0 876 3 1 4 0]
acertou:9752 falhou:248	[8 3 1 1 3 7 933 0 2 0]
Accuracy: 0.975	[1 5 16 7 0 0 0 982 2 15]
	[6 0 2 5 4 1 1 7 944 4]
acertou:9752 falhou:248	[8 5 1 1 8 4 0 9 4 969]]



5. Use the model (1)

```
def make_prediction(model, img):  
    img = img.reshape(1, 1, 28, 28)  
    print(img.shape)  
    print(img.dtype)  
    img = img.to(device) #valves  
    prediction = model(img).cpu().detach().numpy()[0].argmax()  
    print("predict:",prediction)  
    img=img.cpu()  
    plt.axis('off')  
    plt.grid(b=None)  
    plt.imshow(img[0,0], cmap=plt.get_cmap('gray'))  
    plt.show()  
  
model= torch.load('CNNModel_1.pth')  
images, label = next(iter(test_dl))  
make_prediction(model,images[3])
```

5. Use the model (1)

Sequential Layer 1(Conv, ReLU, MaxPool)
Sequential Layer 2(Conv, ReLU, MaxPool)
Linear
ReLU
Linear
Softmax

```
def make_prediction(model, img):  
    img = img.reshape(1, 1, 28, 28)  
    print(img.shape)  
    print(img.dtype)  
    img = img.to(device) #valves  
    prediction = model(img).cpu().detach().numpy()[0].argmax()  
    print("predict:",prediction)  
    img=img.cpu()  
    plt.axis('off')  
    plt.grid(b=None)  
    plt.imshow(img[0,0], cmap=plt.get_cmap('gray'))  
    plt.show()
```

```
model= torch.load('CNNModel1_1.pth')  
images, label = next(iter(test_dl))  
make_prediction(model,images[3])
```

```
torch.Size([1, 1, 28, 28])  
torch.float32  
predict: 4
```



Exercise 4

- Apply the same process to models 2, 3 and 4, improve and present the best value, detailing the best model

Exercise 4

epochs		
batch size		
lesrning rate		
size splits	test: train:	test: train:
layers + activation functions		
loss function		
optimization function		
accuracy		

2. Define the model (2)

```
class CNNModel_2(Module):
    def __init__(self):
        super(CNNModel_2, self).__init__()
        self.layer1 = Sequential(Conv2d(in_channels=1, out_channels=16, kernel_size=3, stride=1, padding=0),
                                   ReLU(),
                                   MaxPool2d(kernel_size=2)
                                   )
        self.layer2 = Sequential(Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=0),
                                   ReLU(),
                                   MaxPool2d(kernel_size=2)
                                   )
        self.fc1 = Linear(32 * 5 * 5, 10)
```

2. Define the model (2)

```
def forward(self, x):  
    out = self.layer1(x)  
    out = self.layer2(out)  
    out = out.view(out.size(0), -1)  
    out = self.fc1(out)  
    return out
```

```
model = CNNModel_2()  
print(summary(model, input_size=(BATCH_SIZE, 1,28,28), verbose=0))
```

2. Define the model (3)

Sequential Layer 1(Conv, BatchNorm, ReLU, MaxPool)
Sequential Layer 2(Conv, BatchNorm, ReLU, MaxPool)
Linear
Dropout
Linear
Linear

```
class CNNModel_3(Module):
    def __init__(self):
        super(CNNModel_3, self).__init__()
        self.layer1 = nn.Sequential(Conv2d(in_channels=1, out_channels=32, kernel_size=3, padding=1),
                                     BatchNorm2d(32),
                                     ReLU(),
                                     MaxPool2d(kernel_size=2, stride=2)
                                    )
        self.layer2 = nn.Sequential(Conv2d(in_channels=32, out_channels=64, kernel_size=3),
                                     BatchNorm2d(64),
                                     ReLU(),
                                     MaxPool2d(2)
                                    )
        self.fc1 = Linear(in_features=64*6*6, out_features=600)
        self.drop = nn.Dropout2d(0.25)
        self.fc2 = Linear(in_features=600, out_features=120)
        self.fc3 = Linear(in_features=120, out_features=10)
```

2. Define the model (3)

Sequential Layer 1(Conv, BatchNorm, ReLU, MaxPool)
Sequential Layer 2(Conv, BatchNorm, ReLU, MaxPool)
Linear
Dropout
Linear
Linear

```
def forward(self, x):  
    out = self.layer1(x)  
    out = self.layer2(out)  
    out = out.view(out.size(0), -1)  
    out = self.fc1(out)  
    out = self.drop(out)  
    out = self.fc2(out)  
    out = self.fc3(out)  
    return out
```

```
model = CNNModel_3()  
print(summary(model, input_size=(BATCH_SIZE, 1,28,28), verbose=0))
```


2. Define the model (4)

```
class CNNModel_4(Module):  
    def __init__(self):  
        super(CNNModel_4, self).__init__()  
        self.layer1 = Sequential(nn.Conv2d(in_channels=1, out_channels=32, kernel_size=5, padding=0),  
                                   BatchNorm2d(32),  
                                   ReLU(),  
                                   MaxPool2d(2),  
                                   Dropout2d(0.2)  
                                   )  
        self.fc1 = Linear(in_features=32*12*12, out_features=128)  
        self.fc2 = Linear(in_features=128, out_features=10)
```

2. Define the model (4)

```
def forward(self, x):  
    out = self.layer1(x)  
    out = out.view(out.size(0), -1)  
    out = self.fc1(out)  
    out = self.fc2(out)  
    return out
```

```
model = CNNModel_4()  
print(summary(model, input_size=(BATCH_SIZE, 1,28,28), verbose=0))
```

Exercise 5

- Apply the same models to fashion-MNIST dataset, improve and present the best value

Exercise 5

epochs		
batch size		
lesrning rate		
size splits	test: train:	test: train:
layers + activation functions		
loss function		
optimization function		
accuracy		

1.1 Visualize the data

```
def output_label(label):
    output_mapping = { 0: "T-shirt/Top", 1: "Trouser", 2: "Pullover",
                      3: "Dress",        4: "Coat",    5: "Sandal",
                      6: "Shirt",        7: "Sneaker", 8: "Bag",
                      9: "Ankle Boot" }
    output_mapping2 = { 0: "0", 1: "1", 2: "2", 3: "3", 4: "4",
                      5: "5", 6: "6", 7: "7", 8: "8", 9: "9"}
    input = (label.item() if type(label) == torch.Tensor else label)
    return output_mapping[input]
```

CNN for multiclass image classification

CIFAR-10 (Canadian Institute For Advanced Research) Dataset

- Image dataset
- Contains 60 000 colour images with 32x32 pixels classified in 10 different classes
- The classes are: planes, cars, birds, cats, deer, frogs, horses, ships and trucks
- There are 6 000 images for each class
- 5 000 images are used for training and 1 000 for testing

0. Prepare the setup

Install pytorch (if needed)

Imports

Constants

```
PATH = './cifar/'  
PATH_CLASSES = './cifar/labels.txt'  
PATH_TRAIN = './cifar/train'  
PATH_TEST = './cifar/test'
```

```
BATCH_SIZE = 128
```

Device management (optional)

1. Prepare the data

```
def get_classes(path):
    with open("cifar/labels.txt") as fich_labels:
        labels = fich_labels.read().split()
        classes = dict(zip(labels, list(range(len(labels)))))
    return classes

dic_classes=get_classes(PATH_CLASSES)
print(dic_classes)

def preprocessor(imagem):
    imagem = np.array(imagem)
    cifar_mean = np.array([0.4914, 0.4822, 0.4465]).reshape(1,1,-1)
    cifar_std = np.array([0.2023, 0.1994, 0.2010]).reshape(1,1,-1)
    imagem = (imagem - cifar_mean) / cifar_std
    xmax, xmin = imagem.max(), imagem.min()
    imagem = (imagem - xmin)/(xmax - xmin)
    imagem = imagem.transpose(2,1,0)
    return imagem
```


1. Prepare the data

```
def get_classes(path):
    with open("cifar/labels.txt") as fich_labels:
        labels = fich_labels.read().split()
        classes = dict(zip(labels, list(range(len(labels)))))
    return classes

dic_classes=get_classes(PATH_CLASSES)
print(dic_classes)           {'airplane': 0, 'automobile': 1, 'bird': 2, 'cat': 3, 'deer': 4, 'dog': 5, 'frog': 6, 'horse': 7, 'ship': 8, 'truck': 9}

def preprocessar(imagem):
    imagem = np.array(imagem)
    cifar_mean = np.array([0.4914, 0.4822, 0.4465]).reshape(1,1,-1)
    cifar_std  = np.array([0.2023, 0.1994, 0.2010]).reshape(1,1,-1)
    imagem = (imagem - cifar_mean) / cifar_std
    xmax, xmin = imagem.max(), imagem.min()
    imagem = (imagem - xmin)/(xmax - xmin)
    imagem = imagem.transpose(2,1,0)
    return imagem
```

1. Prepare the data

```
class Cifar10Dataset(Dataset):
    def __init__(self, path, mun_imagens = 0, transforms=None):
        files = os.listdir(path)
        files = [os.path.join(path,f) for f in files]
        if mun_imagens == 0:
            mun_imagens = len(files)
        self.mun_imagens = mun_imagens
        self.files = random.sample(files, self.mun_imagens)
        self.transforms = transforms

    def __len__(self):
        return self.mun_imagens
```

```
    def __getitem__(self, idx):
        fich_imagem = self.files[idx]
        imagem = Image.open(fich_imagem)
        imagem = preprocessar(imagem)
        label_classe = fich_imagem[:-4].split("_")[-1]
        label = dic_classes[label_classe]
        imagem = imagem.astype(np.float32)
        if self.transforms:
            imagem = self.transforms(imagem)
        return imagem, label
```

1. Prepare the data

```
def prepare_data_loaders(path_train, path_test):
    dataset_train = Cifar10Dataset(path_train, transforms=None)
    dataset_test = Cifar10Dataset(path_test, transforms=None)

    train_size = int(0.8 * len(dataset_train))
    val_size = len(dataset_train) - train_size
    train, validation = random_split(dataset_train, [train_size, val_size],
                                     generator=torch.Generator().manual_seed(42))

    train_dl = DataLoader(train, batch_size=BATCH_SIZE, shuffle=True)
    val_dl = DataLoader(validation, batch_size=BATCH_SIZE, shuffle=True)
    test_dl = DataLoader(dataset_test, batch_size=BATCH_SIZE, shuffle=True)
    train_dl_all = DataLoader(train, batch_size=len(train), shuffle=True)
    val_dl_all = DataLoader(validation, batch_size=len(validation), shuffle=True)
    test_dl_all = DataLoader(dataset_test, batch_size=len(dataset_test), shuffle=True)
    return train_dl, val_dl, test_dl, train_dl_all, val_dl_all, test_dl_all

train_dl, val_dl, test_dl, train_dl_all, val_dl_all, test_dl_all = prepare_data_loaders(PATH_TRAIN, PATH_TEST)
```

1.1 Visualize the data

```
def output_label(label,mapping='label'):
    if mapping == 'ext':
        output_mapping = { 0:"zero", 1:"um", 2:"dois", 3:"tres", 4:"quatro", 5:"cinco", 6:"seis",
7:"sete",
                        8:"oito", 9:"nove" }
    elif mapping == 'ext2':
        output_mapping = { "0":"zero", "1":"um", "2":"dois", "3":"tres", "4":"quatro", "5":"cinco",
                        "6":"seis", "7":"sete", "8":"oito", "9":"nove" }
    else:
        output_mapping = { 0: "0", 1: "1", 2: "2", 3: "3", 4: "4", 5: "5", 6: "6", 7: "7", 8: "8", 9: "9"}
    input = (label.item() if type(label) == torch.Tensor else label)
    return output_mapping[input]
```

1.1 Visualize the data

```
from IPython.display import display
def visualize_data(path):
    ...

def visualize_dataset(train_dl, test_dl, dataset_train, dataset_test):
    ...

visualize_dataset(train_dl, test_dl, train_dl_all, test_dl_all)


def visualize_images(dl):
    ...

visualize_images(train_dl)
```

1.1 Visualize the data

```
Quantidade de casos de Treino:40000
Quantidade de casos de Validação:10000
Quantidade de casos de Teste:10000
Shape tensor batch casos treino, input: torch.Size([128, 3, 32, 32]), output: torch.Size([128])
Shape tensor batch casos validação, input: torch.Size([128, 3, 32, 32]), output: torch.Size([128])
Shape tensor batch casos test, input: torch.Size([128, 3, 32, 32]), output: torch.Size([128])
Valor maximo:1.0 Valor mínimo:0.0
Valor maximo:1.0 Valor mínimo:0.0
tensor([8, 8, 7, 5, 0, 3, 6, 9, 5, 7, 8, 0, 5, 5, 0, 5, 0, 2, 1, 1, 3, 7, 7, 6,
        9, 5, 3, 0, 2, 6, 5, 1, 5, 1, 8, 1, 7, 8, 9, 4, 8, 3, 6, 0, 7, 8, 1, 1,
        8, 6, 5, 0, 7, 4, 6, 6, 3, 4, 9, 6, 6, 3, 4, 5, 5, 6, 2, 1, 1, 2, 5, 2,
        7, 9, 0, 8, 7, 2, 3, 0, 8, 4, 7, 4, 5, 9, 5, 9, 3, 4, 6, 4, 4, 0, 9, 9,
        6, 0, 8, 8, 1, 0, 4, 8, 6, 7, 7, 1, 9, 2, 5, 5, 3, 7, 7, 9, 6, 0, 4, 2,
        8, 2, 6, 0, 8, 3, 4, 6])
```

```
torch.Size([128, 3, 32, 32])
```



1.2 Verify the dataset balancing

2. Define the model (residual)

```
class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1):
        super(ResidualBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=(3, 3),
                                stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(in_channels=out_channels, out_channels=out_channels, kernel_size=(3, 3),
                                stride=1, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.shortcut = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=(1, 1),
                           stride=stride, bias=False),
                nn.BatchNorm2d(out_channels)
            )
```

2. Define the model (residual)

```
def forward(self, x):  
    out = nn.ReLU()(self.bn1(self.conv1(x)))  
    out = self.bn2(self.conv2(out))  
    out += self.shortcut(x)  
    out = nn.ReLU()(out)  
    return out
```

```
class ResNet(nn.Module):  
    def __init__(self, num_classes=10):  
        super(ResNet, self).__init__()  
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=(3, 3), stride=1, padding=1, bias=False)  
        self.bn1 = nn.BatchNorm2d(64)  
        self.block1 = self._create_block(64, 64, stride=1)  
        self.block2 = self._create_block(64, 128, stride=2)  
        self.block3 = self._create_block(128, 256, stride=2)  
        self.block4 = self._create_block(256, 512, stride=2)  
        self.linear = nn.Linear(512, num_classes)
```


2. Define the model (residual)

```
def _create_block(self, in_channels, out_channels, stride):  
    return nn.Sequential(  
        ResidualBlock(in_channels, out_channels, stride),  
        ResidualBlock(out_channels, out_channels, 1)  
    )
```

```
def forward(self, x):  
    out = nn.ReLU()(self.bn1(self.conv1(x)))  
    out = self.block1(out)  
    out = self.block2(out)  
    out = self.block3(out)  
    out = self.block4(out)  
    out = nn.AvgPool2d(4)(out)  
    out = out.view(out.size(0), -1)  
    out = self.linear(out)  
    return out
```

```
model = ResNet()
```

```
print(summary(model, input_size=(BATCH_SIZE, 3,32,32), verbose=0))
```

2. Define the model (residual)

```
Conv Layer 1
Batch Norm 1
Conv Layer 2
Batch Norm 2
Sequential
```

Layer (type:depth-idx)	Output Shape	Param #
Conv2d: 1-1	[128, 64, 32, 32]	1,728
BatchNorm2d: 1-2	[128, 64, 32, 32]	128
Sequential: 1-3	[128, 64, 32, 32]	--
ResidualBlock: 2-1	[128, 64, 32, 32]	--
Conv2d: 3-1	[128, 64, 32, 32]	36,864
BatchNorm2d: 3-2	[128, 64, 32, 32]	128
Conv2d: 3-3	[128, 64, 32, 32]	36,864
BatchNorm2d: 3-4	[128, 64, 32, 32]	128
Sequential: 3-5	[128, 64, 32, 32]	--
ResidualBlock: 2-2	[128, 64, 32, 32]	--
Conv2d: 3-6	[128, 64, 32, 32]	36,864
BatchNorm2d: 3-7	[128, 64, 32, 32]	128
Conv2d: 3-8	[128, 64, 32, 32]	36,864
BatchNorm2d: 3-9	[128, 64, 32, 32]	128
Sequential: 3-10	[128, 64, 32, 32]	--
Sequential: 1-4	[128, 128, 16, 16]	--
ResidualBlock: 2-3	[128, 128, 16, 16]	--
Conv2d: 3-11	[128, 128, 16, 16]	73,728
BatchNorm2d: 3-12	[128, 128, 16, 16]	256
Conv2d: 3-13	[128, 128, 16, 16]	147,456
BatchNorm2d: 3-14	[128, 128, 16, 16]	256
Sequential: 3-15	[128, 128, 16, 16]	8,448
ResidualBlock: 2-4	[128, 128, 16, 16]	--
Conv2d: 3-16	[128, 128, 16, 16]	147,456
BatchNorm2d: 3-17	[128, 128, 16, 16]	256
Conv2d: 3-18	[128, 128, 16, 16]	147,456
BatchNorm2d: 3-19	[128, 128, 16, 16]	256
Sequential: 3-20	[128, 128, 16, 16]	--

```

Sequential: 1-5      [128, 256, 8, 8]      --
├─ResidualBlock: 2-5 [128, 256, 8, 8]      --
│   ├─Conv2d: 3-21    [128, 256, 8, 8]      294,912
│   │   └─BatchNorm2d: 3-22 [128, 256, 8, 8]      512
│   │   └─Conv2d: 3-23    [128, 256, 8, 8]      589,824
│   │   └─BatchNorm2d: 3-24 [128, 256, 8, 8]      512
│   └─Sequential: 3-25 [128, 256, 8, 8]      33,280
├─ResidualBlock: 2-6 [128, 256, 8, 8]      --
│   ├─Conv2d: 3-26    [128, 256, 8, 8]      589,824
│   │   └─BatchNorm2d: 3-27 [128, 256, 8, 8]      512
│   │   └─Conv2d: 3-28    [128, 256, 8, 8]      589,824
│   │   └─BatchNorm2d: 3-29 [128, 256, 8, 8]      512
│   └─Sequential: 3-30 [128, 256, 8, 8]      --
└─Sequential: 1-6    [128, 512, 4, 4]      --
    ├─ResidualBlock: 2-7 [128, 512, 4, 4]      --
    │   ├─Conv2d: 3-31    [128, 512, 4, 4]      1,179,648
    │   │   └─BatchNorm2d: 3-32 [128, 512, 4, 4]      1,024
    │   │   └─Conv2d: 3-33    [128, 512, 4, 4]      2,359,296
    │   │   └─BatchNorm2d: 3-34 [128, 512, 4, 4]      1,024
    │   └─Sequential: 3-35 [128, 512, 4, 4]      132,096
    └─ResidualBlock: 2-8 [128, 512, 4, 4]      --
        ├─Conv2d: 3-36    [128, 512, 4, 4]      2,359,296
        │   └─BatchNorm2d: 3-37 [128, 512, 4, 4]      1,024
        │   └─Conv2d: 3-38    [128, 512, 4, 4]      2,359,296
        │   └─BatchNorm2d: 3-39 [128, 512, 4, 4]      1,024
        └─Sequential: 3-40 [128, 512, 4, 4]      --
└─Linear: 1-7         [128, 10]          5,130

```

```
Total params: 11,173,962
Trainable params: 11,173,962
Non-trainable params: 0
Total mult-adds (T): 1.34
=====
Input size (MB): 1.57
Forward/backward pass size (MB): 1258.30
Params size (MB): 44.70
Estimated Total Size (MB): 1304.57
=====
```

3. Train the model (residual)

51

```
def train_model(h5_file, train_dl, val_dl, model, criterion, optimizer):  
    ...
```

For ResNet model:

```
model = ResNet()  
print(summary(model, input_size=(BATCH_SIZE, 3,32,32), verbose=0))  
EPOCHS = 30  
LEARNING_RATE = 0.001  
criterion = CrossEntropyLoss()  
optimizer = SGD(model.parameters(), lr=LEARNING_RATE)  
starttime = time.perf_counter()  
train_model('CNNModel_cifar_Resnet.pth', train_dl, val_dl, model, criterion, optimizer)  
endtime = time.perf_counter()  
print(f"Tempo gasto: {endtime - starttime} segundos")
```

4. Evaluate the model (residual)

52

```
def evaluate_model(test_dl, model):  
    ...  
def display_predictions(actual_values, predictions):  
    ...  
def display_confusion_matrix(cm,list_classes):  
    ...  
actual_values, predictions = evaluate_model(test_dl_all, model)  
  
model= torch.load('CNNModel_cifar_Resnet.pth')  
actual_values, predictions = evaluate_model(test_dl_all, model)  
display_predictions(actual_values, predictions )  
print(classification_report(actual_values, predictions))  
cr =classification_report(actual_values, predictions, output_dict=True)  
list_classes=[output_label(n,'ext2') for n in list(cr.keys())[0:10] ]  
cm = confusion_matrix(actual_values, predictions)  
print (cm)  
display_confusion_matrix(cm,list_classes)
```

5. Use the model (residual)

```
def make_prediction(model, img):  
    ...  
  
model= torch.load('CNNModel_cifar_Resnet.pth')  
images, label = next(iter(test_dl))  
make_prediction(model,images[3])
```

Exercise 6

- Apply the same process to models 1, 2, 3 and 4, improve and present the best value, detailing the best model

Exercise 6

epochs		
batch size		
lesrning rate		
size splits	test: train:	test: train:
layers + activation functions		
loss function		
optimization function		
accuracy		